


High Performance Computing



Introduction


Levels of Parallelism

Shafaque Fatma Syed


**Assistant Professor - Dept. of Computer Engineering
A P Shah Institute of Technology, Mumbai**

Topics to be discussed

- Levels of Parallelism



Levels of Parallelism




Bit-level Parallelism:

- This form of parallelism is **based on doubling the processor word size**. Increased bit-level parallelism means faster execution of arithmetic operations for large numbers.
- For example, an 8-bit processor needs two cycles to execute a 16-bit addition while a 16-bit processor just one cycle.
- This level of parallelism seems to have come to an end with the introduction of 64-bit processors.

Instruction-level parallelism (ILP):

- This type of parallelism is trying to exploit the potential overlap between instructions in a computer program.
- Multiple instructions from the **same instruction stream** can be executed concurrently.
- Generated and managed by **hardware** (ex:superscalar) or by **compiler** (ex:VLIW).
- Limited in practice by data and control dependencies.

Levels of Parallelism



Instruction-level parallelism (ILP):



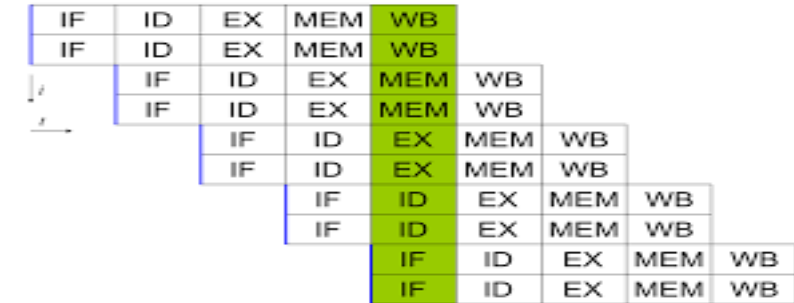
Implemented on hardware level:

Instruction Pipelining — Execute different stages of different independent Instructions in the same cycle thus making full use of idle resources.


Superscalar Processors — Utilize multiple execution units in a single processor die, thus following instructions can be executed without waiting on complex preceding instructions to finish executing.

Out-of-order execution — Instructions not violating any data dependencies are executed when a unit is available even when preceding instructions are still executed.

Speculative execution / Branch prediction — The processor tries to avoid stalling while control instructions (branches — i.e. if or case commands) are still resolved by executing the most probable flow of the program.



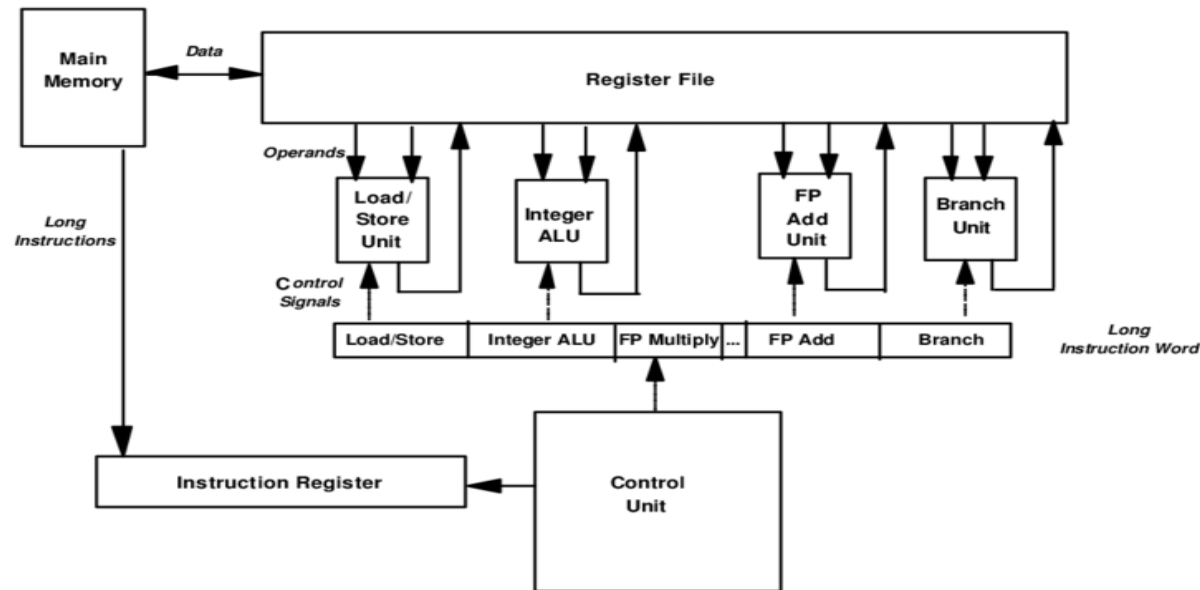
Levels of Parallelism



Instruction-level parallelism (ILP):


Realized on software level:

Using specialized compilers for *Very long instruction word* (VLIW) processors. VLIW processors have multiple execution units organized in multiple pipelines and require the compilers to prepare the parallel instruction streams to fully utilize their resources.



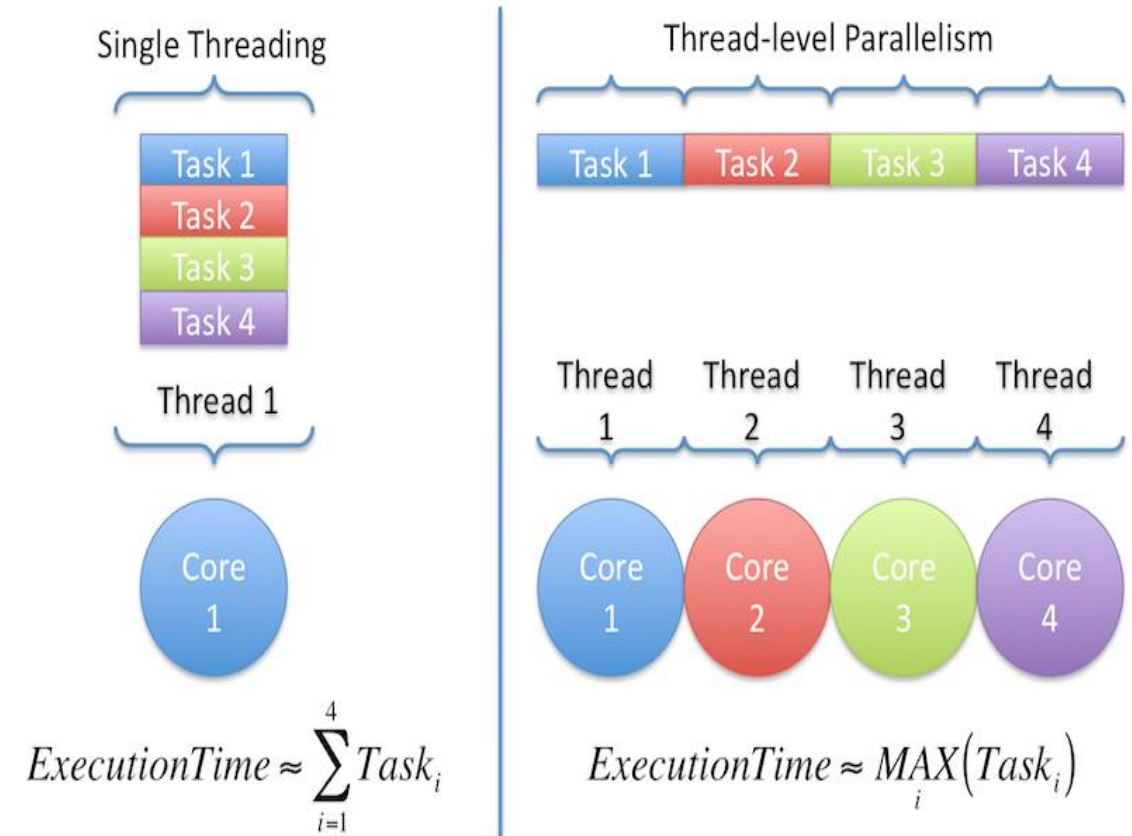
Note: Most processors use a combination of the above ILP techniques. For example, Intel Pentium series used all of the above techniques to achieve higher performance with each product generation.

Levels of Parallelism




Function/ Task / Thread-level parallelism:

- This high-level form of parallel computing is focusing on partitioning the application to be **executed in distinct tasks or threads**, that can be then executed simultaneously on different computation units.
- Threads can work on independent data fragments or even share data between them.
- Multiple threads or instruction sequences from the same application can be executed concurrently
- Generated by compiler and managed by compiler and hardware
- Limited in practice by communication/ synchronization overheads and by algorithm characteristics




Levels of Parallelism



Data parallelism:

- Instructions from a single stream operate concurrently on **several data**
- Limited by non-regular data manipulation patterns and by memory bandwidth
- Tries to partition the data to different available computation units instead.
- The cores execute the *same* task code over the data assigned to each.
- Data parallelism is the only available option for high-level parallelism in computer graphics because the graphic processor units (GPUs) are designed to execute each graphic processing task as fast as possible by partitioning each frame in regions.
- The task on command is then executed independently on each data region by their hundreds of processing units.


Levels of Parallelism



Data Parallelisms	Task Parallelisms
1. Same task are performed on different subsets of same data.	1. Different task are performed on the same or different data.
2. Synchronous computation is performed.	2. Asynchronous computation is performed.
3. As there is only one execution thread operating on all sets of data, so the speedup is more.	3. As each processor will execute a different thread or process on the same or different set of data, so speedup is less.
4. Amount of parallelization is proportional to the input size.	4. Amount of parallelization is proportional to the number of independent tasks is performed.
5. It is designed for optimum load balance on multiprocessor system.	5. Here, load balancing depends upon on the e availability of the hardware and scheduling algorithms like static and dynamic scheduling.


Transaction-level parallelism: Multiple threads/processes from different transactions can be executed concurrently—Limited by concurrency overheads

Memory-level parallelism (MLP): is a term in **computer architecture** referring to the ability to have pending multiple memory operations, in particular cache misses or translation lookaside buffer(TLB) misses, at the same time.



You have completed this topic, you should be able to:

Explain different levels of parallelism?



References Used

- **M. R. Bhujade, —Parallel Computing, 2nd edition, New Age International Publishers, 2009.**