



# High Performance Computing



## HPC Programming


**Shafaque Fatma Syed**  
**Assistant Professor - Dept. of Information  
Technology**  
**A P Shah Institute of Technology, Mumbai**

# Topics to be discussed




- **MPI (Message Passing Interface)**
- **Principles of Message-Passing Programming**
- **The Building Blocks: Send and Receive Operations**
- **MPI: the Message Passing Interface**
- **Topologies and Embedding**
- **Overlapping Communication with Computation**
- **Collective Communication and Computation Operations**

# Let's get started with a small introductory video




<https://www.youtube.com/watch?v=kHV6wmG35po>

# Principles of Message-Passing Programming




- The logical view of a machine supporting the message-passing paradigm consists of  $p$  processes, **each with its own exclusive virtual address space**.
- Each data element must belong to one of the partitions of the space; hence, data must be explicitly partitioned and placed.
- All interactions (read-only or read/write) require cooperation of two processes - the process that has the data and the process that wants to access the data.
- These two constraints make underlying costs very explicit to the programmer.



# Principles of Message-Passing Programming

- Message-passing programs are often written using the *asynchronous or loosely synchronous paradigms*.
- In the asynchronous paradigm, all concurrent tasks execute asynchronously.
- In the loosely synchronous model, tasks or subsets of tasks synchronize to perform interactions. Between these interactions, tasks execute completely asynchronously.
- Most message-passing programs are written using the *single program multiple data (SPMD) model*.



# The Building Blocks: Send and Receive Operations

- The prototypes of these operations are as follows:

**send(void \*sendbuf, int nelems, int dest)**

**receive(void \*recvbuf, int nelems, int source)**

- Consider the following code segments:

**P0**

**a = 100;**

**send(&a, 1, 1);**


**a = 0;**

**P1**

**receive(&a, 1, 0)**


**printf("%d\n", a);**

- The semantics of the send operation require that the value received by process P1 must be 100 as opposed to 0.
- This motivates the design of the send and receive protocols.

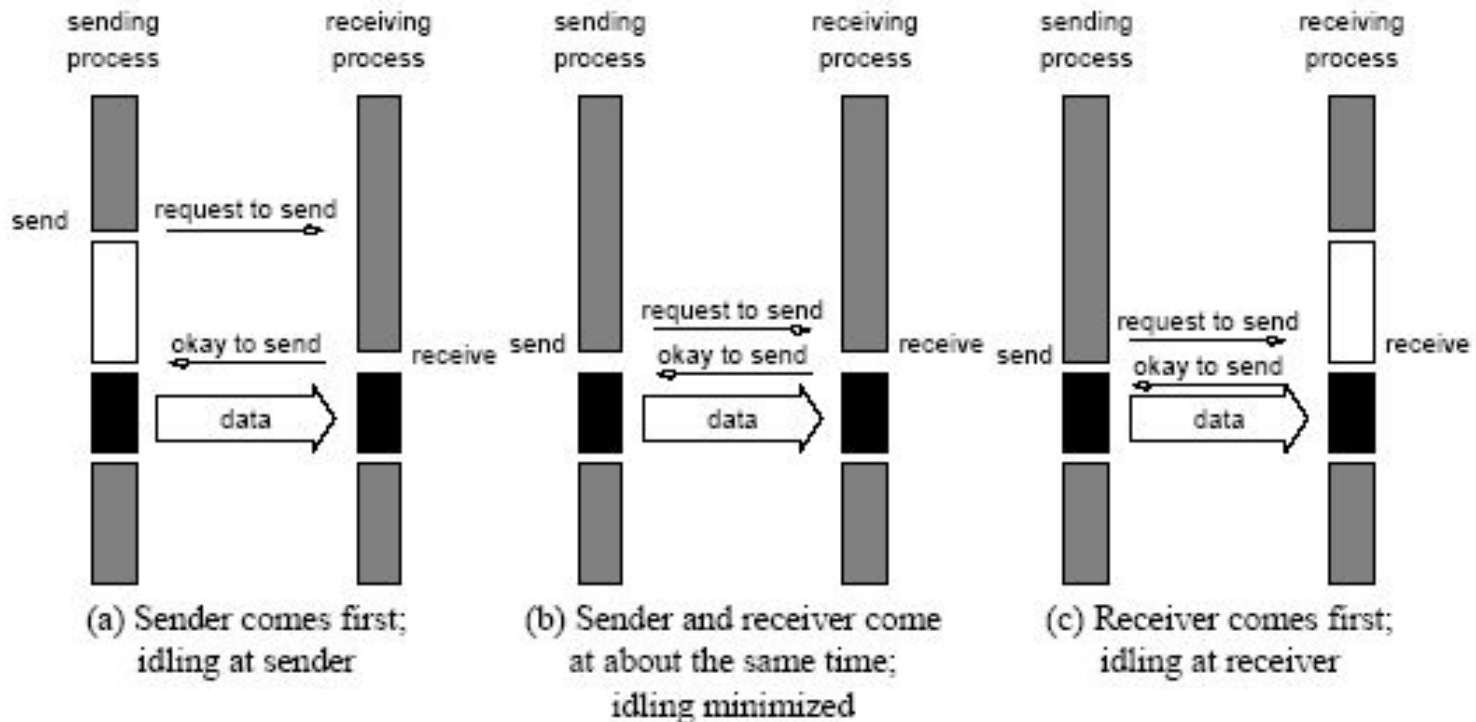


# Non-Buffered Blocking Message Passing Operations

- A simple method for forcing send/receive semantics is for the send operation to **return only when it is safe to do so**.
- In the non-buffered blocking send, the operation **does not return until the matching receive has been encountered** at the receiving process.
- **Idling and deadlocks are major issues** with non-buffered blocking sends.




# Non-Buffered Blocking Message Passing Operations



**Handshake for a blocking non-buffered send/receive operation.**  
It is easy to see that in cases where sender and receiver do not reach communication point at similar times, there can be considerable idling overheads.



# Non-Buffered Blocking Message Passing Operation (Deadlock)




P0

**send (&a, 1, 1)**  
**receive(&b,1,1)**


P1

**send(&a,1,0)**  
**receive(&b,1,0)**




# Buffered Blocking Message Passing Operations


- A simple solution to the idling and deadlocking problem outlined above is to rely on buffers at the sending and receiving ends.
- **In buffered blocking** sends, the sender simply **copies the data into the designated buffer and returns after the copy operation has been completed**. The data is copied at a buffer at the receiving end as well.
- The data must be buffered at the receiving end as well.
- Buffering trades off idling overhead for buffer copying overhead.



# Buffered Blocking Message Passing Operations



**Blocking buffered transfer protocols: (a) in the presence of communication hardware with buffers at send and receive ends; and (b) in the absence of communication hardware, sender interrupts receiver and deposits data in buffer at receiver end.**



# Buffered Blocking Message Passing Operations

Bounded buffer sizes can have significant impact on performance.


P0

```
for (i = 0; i < 1000; i++)  
{  
  produce_data(&a);  
  send(&a, 1, 1);  
}
```

P1

```
for (i = 0; i < 1000; i++)  
{  
  receive(&a, 1, 0);  
  consume_data(&a);  
}
```

**What if consumer was much slower than producer?**



# Buffered Blocking Message Passing Operations

Deadlocks are still possible with buffering since receive operations block.

P0


`receive(&a, 1, 1);`

`send(&b, 1, 1);`

P1


`receive(&a, 1, 0);`

`send(&b, 1, 0);`

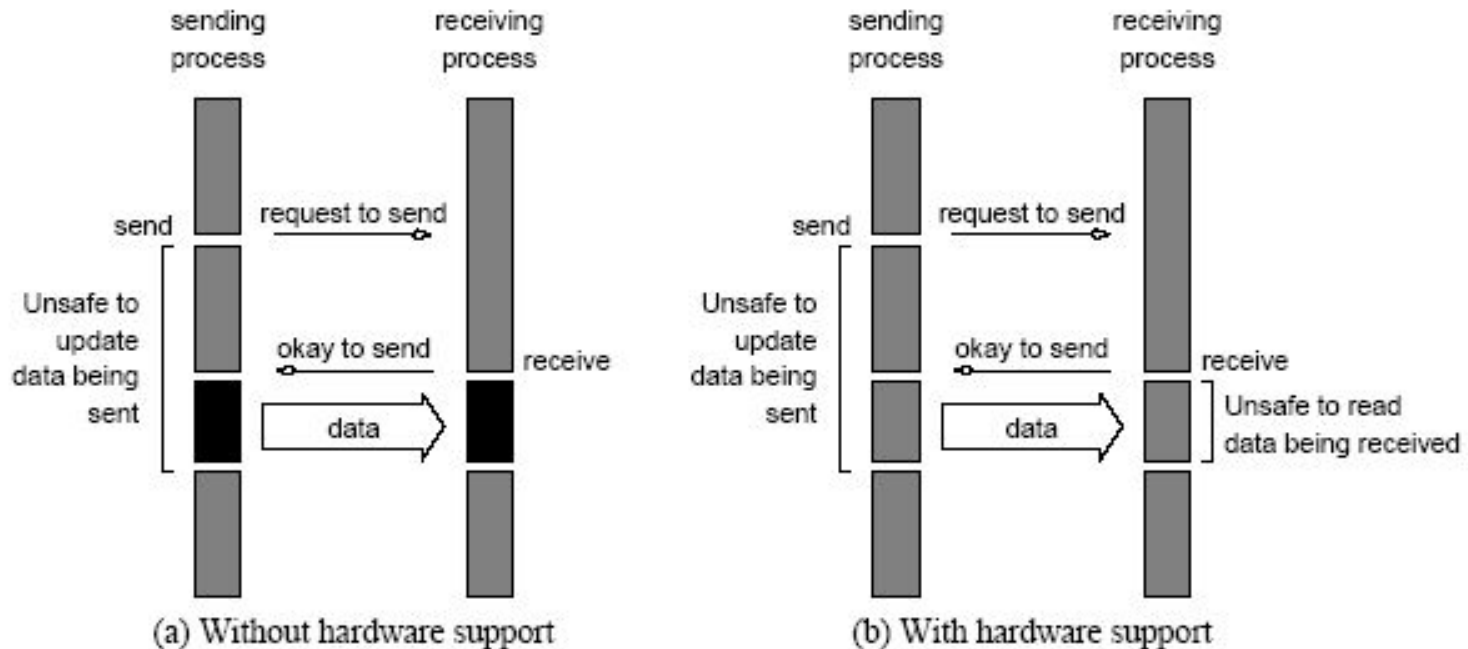


# Non-Blocking Message Passing Operations


- The programmer must ensure semantics of the send and receive.
- This class of non-blocking protocols returns from the send or receive operation before it is semantically safe to do so.
- Non-blocking operations are generally accompanied by a check-status operation.
- When used correctly, these primitives are capable of overlapping communication overheads with useful computations.
- Message passing libraries typically provide both blocking and non-blocking primitives.



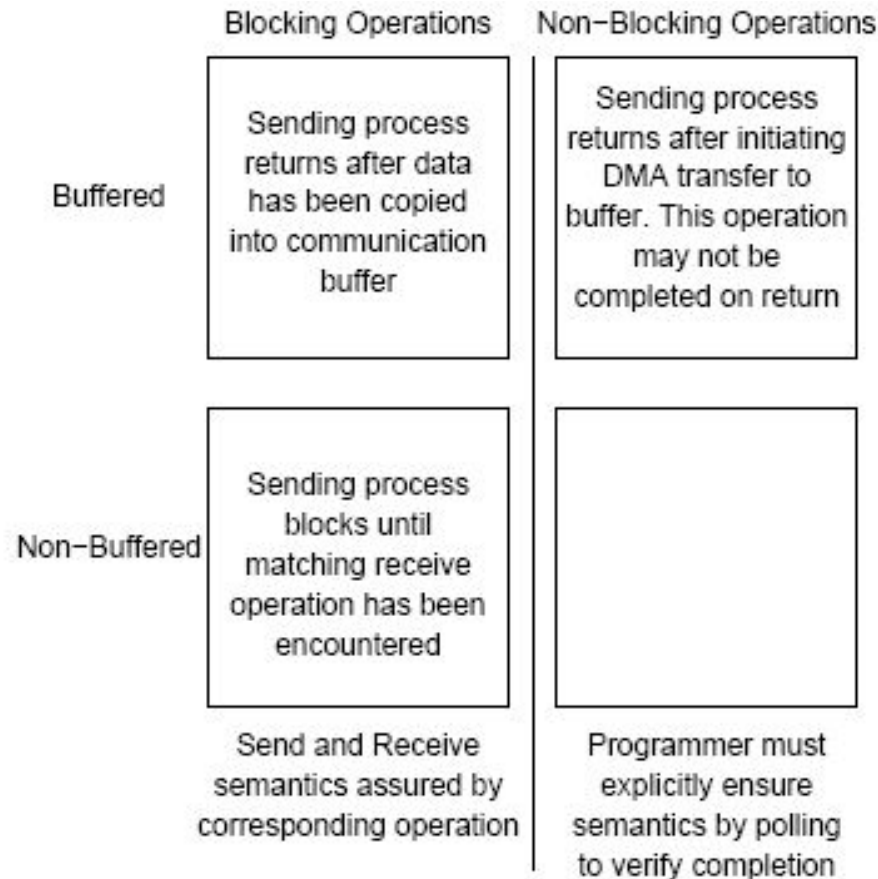
# Non-Blocking Message Passing Operations



Non-blocking non-buffered send and receive operations (a) in absence of communication hardware; (b) in presence of communication hardware.




# Send and Receive Protocols




**Space of possible protocols for send and receive operations.**





# MPI: the Message Passing Interface

- **MPI defines a standard library for message-passing that can be used to develop portable message-passing programs using either C or Fortran.**
- **The MPI standard defines both the syntax as well as the semantics of a core set of library routines.**
- **Vendor implementations of MPI are available on almost all commercial parallel computers.**
- **It is possible to write fully-functional message-passing programs by using only the six routines.**



# MPI: the Message Passing Interface


**The minimal set of MPI routines.**

---

<code>MPI_Init</code>	Initializes MPI.
<code>MPI_Finalize</code>	Terminates MPI.
<code>MPI_Comm_size</code>	Determines the number of processes.
<code>MPI_Comm_rank</code>	Determines the label of calling process.
<code>MPI_Send</code>	Sends a message.
<code>MPI_Recv</code>	Receives a message.

---


# Starting and Terminating the MPI Library




- **MPI\_Init** is called prior to any calls to other MPI routines. Its purpose is to initialize the MPI environment.
- **MPI\_Finalize** is called at the end of the computation, and it performs various clean-up tasks to terminate the MPI environment.
- The prototypes of these two functions are:  


```
int MPI_Init(int *argc, char ***argv)  
int MPI_Finalize()
```
- **MPI\_Init** also strips off any MPI related command-line arguments.
- All MPI routines, data-types, and constants are prefixed by “MPI\_”. The return code for successful completion is **MPI\_SUCCESS**.

# Communicators



- A communicator defines a **communication domain** - a set of processes that are allowed to communicate with each other.
- Information about communication domains is stored in variables of type **MPI\_Comm**.
- Communicators are used as arguments to all message transfer MPI routines.
- A process can belong to many different (possibly overlapping) communication domains.
- MPI defines a default communicator called **MPI\_COMM\_WORLD** which includes all the processes.






# Querying Information

- The **MPI\_Comm\_size** and **MPI\_Comm\_rank** functions are used to determine the number of processes and the label of the calling process, respectively.
- The calling sequences of these routines are as follows:  

```
int MPI_Comm_size(MPI_Comm comm, int *size)  
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```
- The rank of a process is an integer that ranges from zero up to the size of the communicator minus one.




# Our First MPI Program

```
#include <mpi.h>

main(int argc, char *argv[])
{
    int npes, myrank;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &npes);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    printf("From process %d out of %d, Hello
World!\n", myrank, npes);
    MPI_Finalize();
}
```

# Sending and Receiving Messages




- The basic functions for sending and receiving messages in MPI are the **MPI\_Send** and **MPI\_Recv**, respectively.
- The calling sequences of these routines are as follows :

```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

```
int MPI_Recv(void *buf, int count, MPI_Datatype
datatype, int source, int tag, MPI_Comm comm, MPI_Status
*status)
```


- MPI provides equivalent datatypes for all C datatypes. This is done for portability reasons.
- The datatype **MPI\_BYTE** corresponds to a byte (8 bits) and **MPI\_PACKED** corresponds to a collection of data items that has been created by packing non-contiguous data.
- The message-tag can take values ranging from zero up to the MPI defined constant **MPI\_TAG\_UB**.



# MPI Datatypes


MPI Datatype	C Data Type
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	





# Sending and Receiving Messages

- MPI allows **specification of wildcard arguments** for both source and tag.
- If source is set to **MPI\_ANY\_SOURCE**, then any process of the communication domain can be the source of the message.
- If tag is set to **MPI\_ANY\_TAG**, then messages with any tag are accepted.
- On the receive side, the message must be of length equal to or less than the length field specified.
- **MPI\_ERR\_TRUNCATE** in case of message received is larger in length




# Sending and Receiving Messages

- On the receiving end, the **status variable can be used to get information about the MPI\_Recv operation.**
- The corresponding data structure contains:

```
typedef struct MPI_Status {  
    int MPI_SOURCE;  
    int MPI_TAG;  
    int MPI_ERROR; };
```

- The MPI\_Get\_count function returns the precise count of data items received.

```
int MPI_Get_count(MPI_Status *status, MPI_Datatype  
    datatype, int *count)
```




# Avoiding Deadlocks

Consider:

```
int a[10], b[10], myrank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (myrank == 1) {
    MPI_Recv(b, 10, MPI_INT, 0, 2, MPI_COMM_WORLD);
    MPI_Recv(a, 10, MPI_INT, 0, 1, MPI_COMM_WORLD);
}
...
```

If MPI\_Send is blocking, there is a deadlock.




# Avoiding Deadlocks

Consider the following piece of code, in which process  $i$  sends a message to process  $i + 1$  (modulo the number of processes) and receives a message from process  $i - 1$  (modulo the number of processes).

```
int a[10], b[10], npes, myrank;  
MPI_Status status;  
...  
MPI_Comm_size(MPI_COMM_WORLD, &npes);  
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);  
MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,  
MPI_COMM_WORLD);  
MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,  
MPI_COMM_WORLD);  
...
```


Once again, we have a deadlock if `MPI_Send` is blocking.



# Avoiding Deadlocks

We can break the circular wait to avoid deadlocks as follows:

```
int a[10], b[10], npes, myrank;
MPI_Status status;
...
MPI_Comm_size(MPI_COMM_WORLD, &npes);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank%2 == 1) {
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
MPI_COMM_WORLD);
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
MPI_COMM_WORLD);
}
else {
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
MPI_COMM_WORLD);
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
MPI_COMM_WORLD);
}
...
```




# Sending and Receiving Messages Simultaneously

To exchange messages, MPI provides the following function:

```
int MPI_Sendrecv(void *sendbuf, int sendcount,  
MPI_Datatype senddatatype, int dest, int sendtag, void  
*recvbuf, int recvcount, MPI_Datatype recvdatatype, int  
source, int recvtag, MPI_Comm comm, MPI_Status *status)
```


The arguments include arguments to the send and receive functions. If we wish to use the same buffer for both send and receive, we can use:

```
int MPI_Sendrecv_replace(void *buf, int count,  
MPI_Datatype datatype, int dest, int sendtag,  
int source, int recvtag, MPI_Comm comm,  
MPI_Status *status)
```



# Topologies and Embeddings

- MPI allows a programmer to organize processors into logical  $k$ -d meshes.
- The processor ids in `MPI_COMM_WORLD` can be mapped to other communicators (corresponding to higher-dimensional meshes) in many ways.
- The goodness of any such mapping is determined by the interaction pattern of the underlying program and the topology of the machine.
- MPI does not provide the programmer any control over these mappings.



# Topologies and Embeddings

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

(a) Row-major mapping

0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

(b) Column-major mapping

0	3	4	5
1	2	7	6
14	13	8	9
15	12	11	10


(c) Space-filling curve mapping

0	1	3	2
4	5	7	6
12	13	15	14
8	9	11	10

(d) Hypercube mapping

Different ways to map a set of processes to a two-dimensional grid. (a) and (b) show a row- and column-wise mapping of these processes, (c) shows a mapping that follows a space-filling curve (dotted line), and (d) shows a mapping in which neighboring processes are directly connected in a hypercube.





## Row Mapping (Fig. a)

0 (0,0)


row =  $0/4 = 0$

column =  $0\%4 = 0$

12 (3,0)

row =  $12/4 = 3$

column =  $12\%4 = 0$




# Creating and Using Cartesian Topologies

- We can create cartesian topologies using the function:

```
int MPI_Cart_create(MPI_Comm comm_old, int  
ndims, int *dims, int *periods, int reorder,  
MPI_Comm *comm_cart)
```

This function takes the processes in the old communicator and creates a new communicator with `dims` dimensions.

- Each processor can now be identified in this new cartesian topology by a vector of dimension `dims`.



# Creating and Using Cartesian Topologies


- Since sending and receiving messages still require (one-dimensional) ranks, MPI provides routines to convert ranks to cartesian coordinates and vice-versa.

```
int MPI_Cart_coord(MPI_Comm comm_cart, int rank, int  
maxdims, int *coords)
```

```
int MPI_Cart_rank(MPI_Comm comm_cart, int *coords, int  
*rank)
```

- The most common operation on cartesian topologies is a shift. To determine the rank of source and destination of such shifts, MPI provides the following function:

```
int MPI_Cart_shift(MPI_Comm comm_cart, int dir, int s_step,  
int *rank_source, int *rank_dest)
```



# Overlapping Communication with Computation

- In order to overlap communication with computation, MPI provides a pair of functions for performing non-blocking send and receive operations.


```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype,  
int dest, int tag, MPI_Comm comm, MPI_Request *request)  
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype,  
int source, int tag, MPI_Comm comm, MPI_Request *request)
```

- These operations return before the operations have been completed. Function MPI\_Test tests whether or not the non-blocking send or receive operation identified by its request has finished.

```
int MPI_Test(MPI_Request *request, int *flag, MPI_Status  
*status)
```


- MPI\_Wait waits for the operation to complete.

```
int MPI_Wait(MPI_Request *request, MPI_Status *status)
```



Programmer can use this for explicitly freeing space used by various objects in MPI Program.

```
int MPI_Request_free (MPI_Request *request)
```




# Avoiding Deadlocks

Using non-blocking operations remove most deadlocks. Consider:

```
int a[10], b[10], myrank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (myrank == 1) {
    MPI_Recv(b, 10, MPI_INT, 0, 2, &status, MPI_COMM_WORLD);
    MPI_Recv(a, 10, MPI_INT, 0, 1, &status, MPI_COMM_WORLD);
}
...
```


Replacing either the send or the receive operations with non-blocking counterparts fixes this deadlock.

# Collective Communication and Computation Operations





- MPI provides an extensive set of functions for performing common collective communication operations.
- Each of these operations is defined over a group corresponding to the communicator.
- All processors in a communicator must call these operations.

# Collective Communication Operations




The barrier synchronization operation is performed in MPI using:

```
int MPI_Barrier(MPI_Comm  
comm)
```






# Collective Communication Operations




The one-to-all broadcast operation is:

```
int MPI_Bcast(void *buf, int count, MPI_Datatype datatype, int  
              source, MPI_Comm comm)
```




# Collective Communication Operations




The all-to-one reduction operation is:

```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count,  
MPI_Datatype datatype, MPI_Op op, int target, MPI_Comm comm)
```

MPI\_Reduce






# Predefined Reduction Operations

Operation	Meaning	Datatypes
MPI_MAX	Maximum	C integers and floating point
MPI_MIN	Minimum	C integers and floating point
MPI_SUM	Sum	C integers and floating point
MPI_PROD	Product	C integers and floating point
MPI_LAND	Logical AND	C integers
MPI_BAND	Bit-wise AND	C integers and byte
MPI_LOR	Logical OR	C integers
MPI_BOR	Bit-wise OR	C integers and byte
MPI_LXOR	Logical XOR	C integers
MPI_BXOR	Bit-wise XOR	C integers and byte
MPI_MAXLOC	max-min value-location	Data-pairs
MPI_MINLOC	min-min value-location	Data-pairs

# Collective Communication Operations




- The operation `MPI_MAXLOC` combines pairs of values ( $v_i, l_i$ ) and returns the pair ( $v, l$ ) such that  $v$  is the maximum among all  $v_i$  's and  $l$  is the corresponding  $l_i$  (if there are more than one, it is the smallest among all these  $l_i$  's).
- `MPI_MINLOC` does the same, except for minimum value of  $v_i$ .

Value	15	17	11	12	17	11
Process	0	1	2	3	4	5

`MinLoc(Value, Process) = (11, 2)`

`MaxLoc(Value, Process) = (17, 1)`


An example use of the `MPI_MINLOC` and `MPI_MAXLOC` operators.



# Collective Communication Operations

MPI datatypes for data-pairs used with the `MPI_MAXLOC` and `MPI_MINLOC` reduction operations.

MPI Datatype	C Datatype
<code>MPI_INT</code>	pair of ints
<code>MPI_SHORT_INT</code>	short and int
<code>MPI_LONG_INT</code>	long and int
<code>MPI_LONG_DOUBLE_INT</code>	long double and int
<code>MPI_FLOAT_INT</code>	float and int
<code>MPI_DOUBLE_INT</code>	double and int




# Collective Communication Operations


- If the result of the reduction operation is needed by all processes, MPI provides:

```
int MPI_Allreduce(void *sendbuf, void *recvbuf, int count,  
MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

MPI\_Allreduce

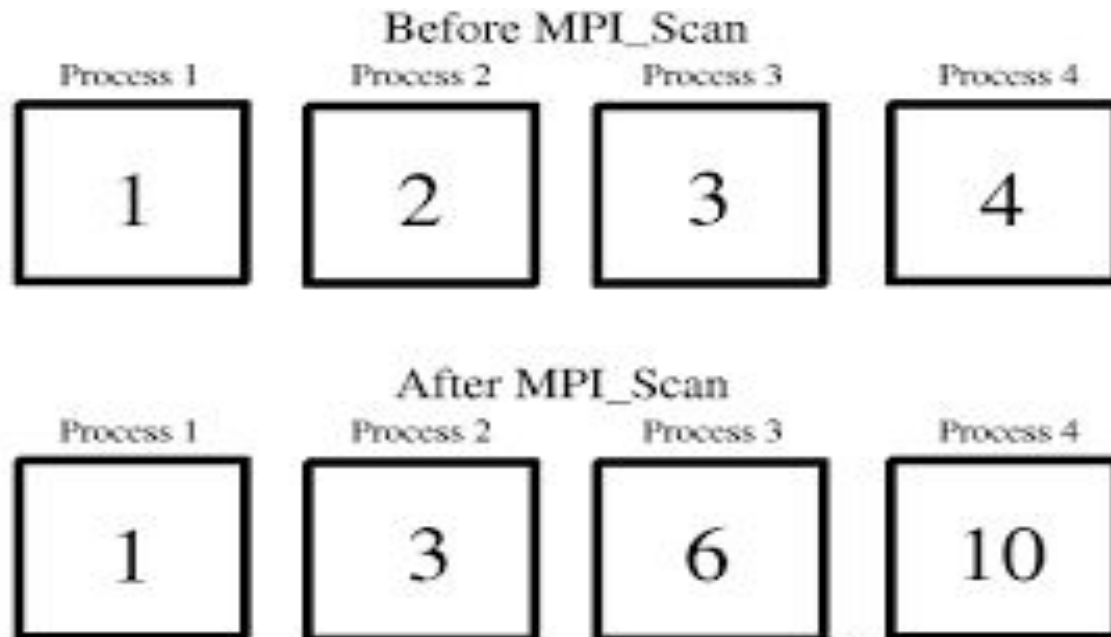


# Collective Communication Operations




- The fig. below shows computation of prefix-sums.

```
int MPI_Scan(void *sendbuf, void *recvbuf, int count,  
MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```




# Collective Communication Operations




- The gather operation is performed in MPI using:

```
int MPI_Gather(void *sendbuf, int sendcount, MPI_Datatype  
senddatatype, void *recvbuf, int recvcount, MPI_Datatype  
recvdatatype, int target, MPI_Comm comm)
```






# Collective Communication Operations




- MPI also provides the MPI\_Allgather function in which the data are gathered at all the processes.

```
int MPI_Allgather(void *sendbuf, int sendcount,  
MPI_Datatype senddatatype, void *recvbuf, int recvcount,  
MPI_Datatype recvdatatype, MPI_Comm comm)
```





# Collective Communication Operations



- The corresponding scatter operation is:

```
int MPI_Scatter(void *sendbuf, int sendcount, MPI_Datatype  
senddatatype, void *recvbuf, int recvcount, MPI_Datatype  
recvdatatype, int source, MPI_Comm comm)
```






# Collective Communication Operations

- The all-to-all personalized communication operation is performed by:

```
int MPI_Alltoall(void *sendbuf, int sendcount,  
MPI_Datatype senddatatype, void *recvbuf, int  
recvcount, MPI_Datatype recvdatatype, MPI_Comm comm)
```

- Using this core set of collective operations, a number of programs can be greatly simplified.




# Collective Communication Operations

Suppose there are four processes including the root, each with arrays as shown below on the left. After the all-to-all operation

```
MPI_Alltoall(u, 2, MPI_INT, v, 2, MPI_INT, MPI_COMM_WORLD);
```

the data will be distributed as shown below on the right:

array u	Rank	array v																
<table><tr><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td><td>16</td><td>17</td></tr></table>	10	11	12	13	14	15	16	17	0	<table><tr><td>10</td><td>11</td><td>20</td><td>21</td><td>30</td><td>31</td><td>40</td><td>41</td></tr></table>	10	11	20	21	30	31	40	41
10	11	12	13	14	15	16	17											
10	11	20	21	30	31	40	41											
<table><tr><td>20</td><td>21</td><td>22</td><td>23</td><td>24</td><td>25</td><td>26</td><td>27</td></tr></table>	20	21	22	23	24	25	26	27	1	<table><tr><td>12</td><td>13</td><td>22</td><td>23</td><td>32</td><td>33</td><td>42</td><td>43</td></tr></table>	12	13	22	23	32	33	42	43
20	21	22	23	24	25	26	27											
12	13	22	23	32	33	42	43											
<table><tr><td>30</td><td>31</td><td>32</td><td>33</td><td>34</td><td>35</td><td>36</td><td>37</td></tr></table>	30	31	32	33	34	35	36	37	2	<table><tr><td>14</td><td>15</td><td>24</td><td>25</td><td>34</td><td>35</td><td>44</td><td>45</td></tr></table>	14	15	24	25	34	35	44	45
30	31	32	33	34	35	36	37											
14	15	24	25	34	35	44	45											
<table><tr><td>40</td><td>41</td><td>42</td><td>43</td><td>44</td><td>45</td><td>46</td><td>47</td></tr></table>	40	41	42	43	44	45	46	47	3	<table><tr><td>16</td><td>17</td><td>26</td><td>27</td><td>36</td><td>37</td><td>46</td><td>47</td></tr></table>	16	17	26	27	36	37	46	47
40	41	42	43	44	45	46	47											
16	17	26	27	36	37	46	47											




# Groups and Communicators


- In many parallel algorithms, communication operations need to be restricted to certain subsets of processes.
- MPI provides mechanisms for partitioning the group of processes that belong to a communicator into subgroups each corresponding to a different communicator.
- The simplest such mechanism is:

```
int MPI_Comm_split(MPI_Comm comm, int color, int key,  
                  MPI_Comm *newcomm)
```

- This operation groups processors by color and sorts resulting groups on the key.




# Groups and Communicators



Using `MPI_Comm_split` to split a group of processes in a communicator into subgroups.

For any one color, the key values do not have to be unique. The `MPI_Comm_split` function sorts processes in order according to the value of the *key* parameter, and it sorts ties by their relative rank in the source group. If the same value is specified for all the *key* parameters, then all the processes in a given color have the same relative rank order that they had in their parent group.




# Groups and Communicators

- In many parallel algorithms, processes are arranged in a virtual grid, and in different steps of the algorithm, communication needs to be restricted to a different subset of the grid.
- MPI provides a convenient way to partition a Cartesian topology to form lower-dimensional grids:


```
int MPI_Cart_sub(MPI_Comm comm_cart, int *keep_dims,  
                MPI_Comm *comm_subcart)
```

- If `keep_dims[i]` is true (non-zero value in C) then the `i`th dimension is retained in the new sub-topology.
- The coordinate of a process in a sub-topology created by `MPI_Cart_sub` can be obtained from its coordinate in the original topology by disregarding the coordinates that correspond to the dimensions that were not retained.




# Groups and Communicators

**Original Topology  
(2 x 4 x 7)**



`keep_dims[] = {true, false, true}`



(a)


**`keepdims[true,  
false,true]`**

**Original topology  
is split into 4  
two-dimensional  
sub-topologies of  
size 2x7**

**`keepdims[false,  
false,true]`**

**Original topology  
is split into 8  
one-dimensional  
sub-topologies of  
size 1x7**


`keep_dims[] = {false, false, true}`



(b)

Splitting a Cartesian topology of size 2 x 4 x 7 into (a) four subgroups of size 2 x 1 x 7, and (b) eight subgroups of size 1 x 1 x 7.





# References Used

- **Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar , —Introduction to Parallel Computing,Pearson Education, Second Edition, 2007.**
- **<https://mpitutorial.com/tutorials/>**